

CS 4310 Algorithms: IV. Greedy Methods

E. de Doncker¹

¹Dept. of Computer Science, WMU, U.S.A.

February 2024

Outline

- 1 Greedy meta-algorithm
- 2 Greedy knapsack
- 3 Optimal merge patterns
- 4 Huffman trees and codes
- 5 Graphs - Terminology
- 6 Minimum spanning tree

Outline

- 1 Greedy meta-algorithm
- 2 Greedy knapsack
- 3 Optimal merge patterns
- 4 Huffman trees and codes
- 5 Graphs - Terminology
- 6 Minimum spanning tree

Outline

- 1 Greedy meta-algorithm
- 2 Greedy knapsack
- 3 Optimal merge patterns
- 4 Huffman trees and codes
- 5 Graphs - Terminology
- 6 Minimum spanning tree

Outline

- 1 Greedy meta-algorithm
- 2 Greedy knapsack
- 3 Optimal merge patterns
- 4 Huffman trees and codes
- 5 Graphs - Terminology
- 6 Minimum spanning tree

Outline

- 1 Greedy meta-algorithm
- 2 Greedy knapsack
- 3 Optimal merge patterns
- 4 Huffman trees and codes
- 5 Graphs - Terminology
- 6 Minimum spanning tree

Outline

- 1 Greedy meta-algorithm
- 2 Greedy knapsack
- 3 Optimal merge patterns
- 4 Huffman trees and codes
- 5 Graphs - Terminology
- 6 Minimum spanning tree

Greedy meta-algorithm

A typical greedy algorithm **optimizes (minimizes or maximizes)** an **objective function** (representing a **cost** or a **gain**) under a set of **constraints**. A solution that satisfies the constraints is called a **feasible solution**. An **optimal solution** is a feasible solution that optimizes the objective function.

A general “**meta**–” or umbrella greedy algorithm is given below [1].

Greedy meta-algorithm

```
Meta-algorithm Greedy(A, n) { // returns solution for inputs A and size n  
  solution =  $\emptyset$ ; // Initialize solution  
  for (i = 1; i ≤ n; i++) { // Construct solution  
    x = Choose(A);  
    if (Feasible(solution, x)) // Check if x leads to feasible solution  
      solution = Union(solution, x); // Add x to solution  
  }  
  return solution;  
}
```


Greedy knapsack problem

Given the knapsack capacity M , number of items n , with p_i and w_i the (positive) profit and weight, respectively, per unit of item i , the problem is to determine the fractions x_i of the items to fill the knapsack so that the total profit is maximized. That is [1],

Greedy knapsack problem

$$\text{maximize } \sum_{i=1}^n p_i x_i \text{ \textit{objective function}, subject to the constraints } \sum_{i=1}^n w_i x_i \leq M$$

and $0 \leq x_i \leq 1$, for $1 \leq i \leq n$

Example : $M = 20$, $n = 3$, profits $\mathbf{p} = (50, 30, 20)$, weights $\mathbf{w} = (10, 15, 5)$.

Strategy: add items to the knapsack in non-increasing order of profit per (unit of) weight. The ratios for the example are: $(50/10, 30/15, 20/5) = (5, 2, 4)$. The first and third object are put in completely, leaving room for one third of the second object. The total profit is $50 + 20 + 30/3 = 80$.

Algorithm *GreedyKnapsack*() [1]

```
Algorithm GreedyKnapsack(  $M, p, w, n, x$  ) { // returns optimal fractions  $x$   
// for inputs  $M$  (knapsack capacity),  $p$  (profits),  $w$  (weights), and size  $n$ ,  
// assuming  $p[1]/w[1] \geq p[2]/w[2] \geq \dots \geq p[n]/w[n]$   
// for positive profits and weights.  
  for (  $i = 1; i \leq n; i++$  )  $x[i] = 0$ ; // Initialize  
   $U = M$ ; //  $U$  is remaining capacity  
  for (  $i = 1; i \leq n; i++$  ) {  
    if (  $w[i] > U$  ) break;  
     $x[i] = 1$ ;  
     $U -= w[i]$ ;  
  }  
  if (  $i \leq n$  )  $x[i] = U/w[i]$ ;  
}
```

Not including sorting of the ratios, the time complexity is linear in n .

Including sorting of the ratios, the time complexity is thus $\mathcal{O}(n \log n)$.

It can be proved that the strategy of adding items to the knapsack in non-increasing order of the profit by unit of weight ratios, yields an optimal solution [1].

Optimal merge patterns


The objective is to **find an optimal merge pattern** for multiple sorted files that need to be merged, using a minimum number of record moves. A two-way binary merge tree is constructed.

As an example, consider files x_1, x_2, x_3, x_4, x_5 of sizes 9, 44, 25, 93, 13. If merges are done by: (1) merging x_1, x_2 using up to $9 + 44 = 53$ moves; (2) merging the previous result with x_3 in $53 + 25 = 78$ moves; (3) merging the previous result with x_4 in $78 + 93 = 171$ moves; (4) merging the previous result with x_5 in $171 + 13 = 184$ moves; then the **total number of moves is $53 + 78 + 171 + 184 = 486$** .

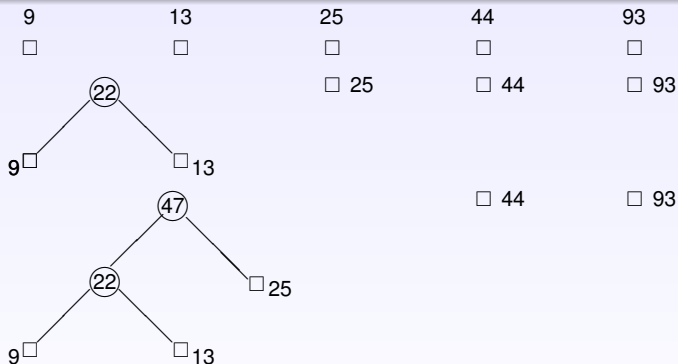
The optimal strategy is to combine the least two at each step: (1) merge x_1 and x_5 in $9 + 13 = 22$ moves, (2) merge the result of length 22 with x_3 in $22 + 25 = 47$ moves, (3) merge the result of length 47 with x_2 in $44 + 47 = 91$ moves, (4) merge the result of length 91 with x_4 in $91 + 93 = 184$ moves; then the **total number of moves is $22 + 47 + 91 + 184 = 344$** .

Optimal merge patterns - Example

The given data are represented by the leaf (square) nodes, called **external nodes**. The nodes combining children nodes are **internal nodes**, labeled by the sum of the keys of the children nodes. The key value (weight) for the i -th external node is q_i and its distance from the root is d_i . The **weighted external path length** is defined as $\sum_{i=1}^n q_i d_i$. We have $\sum_{i=1}^n q_i d_i = (9 + 13) \times 4 + 25 \times 3 + 44 \times 2 + 93 \times 1 = 344$.



Optimal merge patterns - Example



...

Optimal merge patterns - Algorithm

An outline of the algorithm can be given as follows (see also [2, 1]):

- Initialize the tree by forming a list of n one-node binary trees with weights q_i , $i = 1, \dots, n$.
- Repeat until only one tree remains:
 - Find two subtrees of minimum weight
 - Merge into one subtree (create a new parent node and attach the selected subtrees as left and right child (in any order))
 - Set the new node weight to the sum of the weights of the children
 - Insert new subtree into the list

Pseudocode of the algorithm is given in [1]. It is shown that the algorithm generates the **optimal two-way merge tree**. The strategy **minimizes the weighted external path length** $\sum_{i=1}^n q_i d_i$, where d_i is the distance of the external node with weight q_i .

Indeed, this corresponds to placing external nodes with smaller weights farther from the root (at larger depths d_i). Using a **minheap** as a **priority queue data structure** for selecting the next tree with least weight efficiently, the time as a function of n is $\mathcal{O}(n \log n)$. Using a **linked list** requires total time $\mathcal{O}(n^2)$.

Huffman trees and codes

The general objective is **file compression**, in order to minimize the file size and thus the transmission time for sending it. For compressing a text, or sending a message in the English language, all alphabet symbols and punctuation symbols in English need to be encoded. It is thus beneficial to give **shorter encodings to symbols that occur more frequently**.

A **Huffman tree** is constructed using an algorithm as outlined for optimal merge patterns, where the weights q_i represent frequencies. The basic principle is that **more frequent symbols occur at shorter distances from the root**, and **less frequent symbols are deeper in the tree**. The algorithm minimizes the weighted external path length $\sum_{i=1}^n q_i d_i$ and produces an optimal encoding tree.

Using a **heap priority queue** the algorithm time is $\mathcal{O}(n \log n)$.

To construct the **Huffman code as a binary string** for each symbol, descend the tree from the root to the symbol, while assigning '1' to each right branch and '0' to each left branch.

An example from [2] is based on the **relative frequencies of the letters** in the English alphabet (occurring in a text), as given in the table below. Its decoding tree (Huffman tree) is also given.

Huffman trees and codes - Example

Letter	Frequency	Huffman code
A	0.073	1011
B	0.009	000100
C	0.030	01011
D	0.044	0000
E	0.130	100
F	0.028	01010
G	0.016	011110
H	0.035	10100
I	0.074	1100
J	0.002	011100001
K	0.003	01110010
L	0.035	10101
M	0.025	00011

Letter	Frequency	Huffman code
N	0.078	1111
O	0.074	1101
P	0.027	01000
Q	0.003	01110001
R	0.077	1110
S	0.063	0110
T	0.093	001
U	0.027	01001
V	0.013	000101
W	0.016	011101
X	0.005	01110011
Y	0.019	011111
Z	0.001	011100000

Figure: Frequencies and Huffman encoding of English alphabet symbols [2]

Huffman trees and codes - Example




Figure: Huffman (encoding) tree for English alphabet symbols [2]

Graphs - Terminology

Graph $G = (V, E)$: has **set of vertices** V and **set of edges** E .

G can be a **directed** graph (with directed edges), or an **undirected** graph (with undirected edges).

The presence of an edge (v_i, v_j) indicates that vertex v_j is **adjacent** to v_i (vertex v_i is adjacent from v_j ; for an undirected graph, v_i and v_j are just called adjacent).

A **weighted graph** has weights on the edges, weight w_{ij} on edge (v_i, v_j) .

A **path** in G consists of a sequence of vertices v_1, v_2, \dots, v_k (for $k > 1$), such that (v_i, v_{i+1}) is an edge in G for $1 \leq i \leq k - 1$. This path is of length $k - 1$. A path v_1, v_2, \dots, v_k where $v_k = v_1$ is a **cycle**.




Figure: Undirected graph - Cities example [2]

Graphs - Terminology

Graph representations: (i) **adjacency matrix**; (ii) **adjacency list**

(i) For a graph $G = (V, E)$ with n vertices ($|V| = n$), the adjacency matrix is an $n \times n$ **matrix**. For a weighted graph, the (i, j) matrix element is the weight w_{ij} on edge (v_i, v_j) . The adjacency matrix is **symmetric** for an **undirected graph** (see below for the Cities example [2]).

	1	2	3	4	5	6	7	8	9	
1		345					97	230	39	Baltimore
2	345			186	252	445	365	217		Buffalo
3				244	265			284	492	Cincinnati
4		186	244		167	507		125		Cleveland
5		252	265	167						Detroit
6		445		507			92	386		New York
7	97	365				92		305		Philadelphia
8	230	217	284	125		386	305		231	Pittsburgh
9	39		492					231		Washington

Figure: Adjacency matrix representation - Cities example [2]

Graphs - Terminology

(ii) For graph $G = (V, E)$ with n vertices and e edges ($|V| = n$, $|E| = e$), the adjacency list data structure has **n head nodes** (with pointers to n lists). The list for head node i (representing vertex v_i) contains a list node for each vertex v_j adjacent to v_i . There are **e list nodes** for a **directed graph**, and **$2e$ list nodes** for an **undirected graph**, see figure for the Cities example. In the list for head node i , the list node for adjacent vertex v_j has a **field** for the **vertex number (j)**, the **weight w_{ij}** of edge (v_i, v_j) , and a **next pointer**. The list nodes are in no particular order.




Figure: Adjacency list representation - Cities example [2]

Spanning tree

A **spanning tree** of an undirected graph $G = (V, E)$ with $|V| = n$ is a tree formed by the n vertices as its nodes and $n - 1$ edges from E .

A connected graph G may have multiple spanning trees. For a weighted graph G , the **cost** of a spanning tree is the sum of the weights on its edges.

The **minimum cost spanning tree** is a spanning tree of minimum cost.

The figure below [1] shows a graph (left) and minimum cost spanning tree (right).




Figure: Graph and minimum cost spanning tree [1]

Kruskal's algorithm - construction

Kruskal's algorithm for minimum cost spanning tree construction is illustrated for the example [1].




Figure: Kruskal's algorithm for minimum cost spanning tree - construction [1]

Prim's algorithm - construction

Prim's algorithm for minimum cost spanning tree construction is illustrated for the example [1].




Figure: Prim's algorithm for minimum cost spanning tree - construction [1]

Kruskal's algorithm V1 [1]

```
Algorithm Kruskal1 ( $G, n, T$ ) { // returns minimum cost spanning tree  $T$  for graph  $G$ 
1    $T = \emptyset$ ; // Initialize  $T$  empty
2   while (( $T$  has less than  $n - 1$  edges) and ( $E \neq \emptyset$ )) { // Consider next edge to add
3       Choose edge  $(v, w)$  of lowest cost;
4       Delete  $(v, w)$  from  $E$ ;
5       if (( $v, w$ ) does not create a cycle in  $T$ )
6           add  $(v, w)$  to  $T$ ;
7       else discard  $(v, w)$ ;
8   }
```


Kruskal's algorithm V2

```
Algorithm Kruskal2 ( $G, n, T$ ) { // returns minimum cost spanning tree  $T$  for graph  $G$ 
1    $T = \emptyset$ ; // Initialize  $T$  empty (no edges,  $n$  one-node subtrees)
    build heap keyed with edge costs;
2   while (( $T$  has less than  $n - 1$  edges) and ( $E \neq \emptyset$ )) { // Consider next edge to add
3       Choose edge  $(v, w)$  of lowest cost; // delete from heap
4       Delete  $(v, w)$  from  $E$ ;
5       if ( $(v, w)$  does not create a cycle in  $T$ )
          // if  $v$  and  $w$  belong to different subtrees
6           add  $(v, w)$  to  $T$ ; // join (union) the subtrees of  $v$  and  $w$ 
7       else discard  $(v, w)$ ;
8   }
}
```

Kruskal's algorithm - Analysis

Using a **min-heap** keyed with the edge costs:

- Building the heap initially requires $\mathcal{O}(e)$ time, where $e = |E|$.
- Within the while loop, with $\mathcal{O}(e)$ iterations:
 - delete from heap is done in $\mathcal{O}(\log e)$ time,
 - thus $\mathcal{O}(e \log e)$ time throughout the loop.

The accumulated time of lines 5 and 6 throughout the while loop is dominated by the heap processing time by using a suitable data structure for checking whether edge (u, v) would create a cycle.

The **union-find** data structure is used to perform the following efficiently:

```
 $T_1$  = subtree containing  $v$ ; // find( $v$ ) in union-find data structure  
 $T_2$  = subtree containing  $w$ ; // find( $w$ ) in union-find data structure  
if ( $T_1 \neq T_2$ ) // ( $(v, w)$  does not create a cycle in  $T$ )  
    union( $T_1, T_2$ ); // (add  $(v, w)$  to  $T$ )
```

Then the time complexity of Kruskal's algorithm is $\mathcal{O}(e \log e)$.

Prim's algorithm - Analysis

The time complexity of **Prim's algorithm** as given in [1] is $\mathcal{O}(n^2)$; it is stated that using a red-black tree for the set of nodes that have not yet been included in the tree, yields a version of time complexity $\mathcal{O}((n + e) \log n)$.

BIBLIOGRAPHY



S. Sahni E. Horowitz and S. Rajasekeran.

Computer Algorithms/C++.

Computer Science Press, 2nd edition, 1998.

ISBN 0-7167-8315-0.



B. M. E. Moret and H. D. Shapiro.

Algorithms from P to NP, Vol. I - Design and Efficiency.

Benjamin/Cummings Publishing Company, Inc., 1990.

ISBN 0-8053-8008-6.